


Difference Engine: Harnessing Memory Redundancy in Virtual Machines (D. Gupta et al)

Presented by: Konrad Gołuchowski

What is Virtual machine monitor (VMM)?



A few samples of VMM



Why use virtualization in business?

- ▶ Low CPU utilization by individual services (e.g. 5-10%)
- ▶ Need isolation between services
 - ▶ (services may require different configuration)
- ▶ Deploy a new server within minutes
- ▶ And your competitors are already using it!



Source: <http://brandingbrand.com/>



Who uses VMM?

- ▶ Anyone?
- ▶ Any problems with it?
- ▶ Maybe performance problems?



Source: www.jtgraphic.net



Motivation for this paper

- ▶ Server configuration (students):
 - ▶ 78 GB RAM
 - ▶ 24 CPU cores (Intel Xeon @ 2.66GHz)
- ▶ How many virtual machines may work on this server?
 - ▶ 100?
 - ▶ 200?
- ▶ But, even with 100 machines, each machine could use less than 1 GB RAM



Motivation for this paper


- ▶ Hardware upgrade is complicated operation
- ▶ High-capacity memory chips are expensive
- ▶ And they consume power
- ▶ **And... competitors are already doing it!**




Source: <http://news-libraries.mit.edu/>

Competitors?

- ▶ VMware ESX server:
 - ▶ Content-based page sharing
 - ▶ Reduction of memory footprint by 10-40% (homogenous systems)




Proposed solution: Difference Engine



Proposed solution: Difference Engine

- ▶ Built on the base of Xen
- ▶ Main ideas:
 - ▶ Page sharing (as VMware does)
 - ▶ Sub-page level sharing (patching)
 - ▶ In-memory compression




Page Sharing

- ▶ How to locate identical pages quickly?
 - ▶ Use hashing and byte-by-byte comparison
- ▶ Then, we just update virtual memory to point at the shared copy
- ▶ But what if one virtual machine changes shared page?
 - ▶ Mark shared copy as read-only
 - ▶ Writing to the page causes page fault trapped by VMM
 - ▶ VMM creates a private copy and updates virtual memory



Handling similar pages

- ▶ Store a reference page and a patch
- ▶ How to detect similar pages?
 - ▶ Hash 64-byte blocks at random locations
 - ▶ Compare computed hashes
- ▶ Generate patch (or patches and choose best one)
- ▶ Don't use too large patches
- ▶ When use patching?



In-memory compression

- ▶ Basic idea: compress pages that are not similar to anything
- ▶ Compress when it is worth
- ▶ Use compression and patching only for pages accessed infrequently
- ▶ How to locate those infrequently accessed pages?
 - ▶ Not-recently-used policy and global scanning every some time
 - ▶ Scans checks and clears referenced and modified bits
 - ▶ VM have time to reset those bits
 - ▶ Only a part of memory is scanned (for each VM)



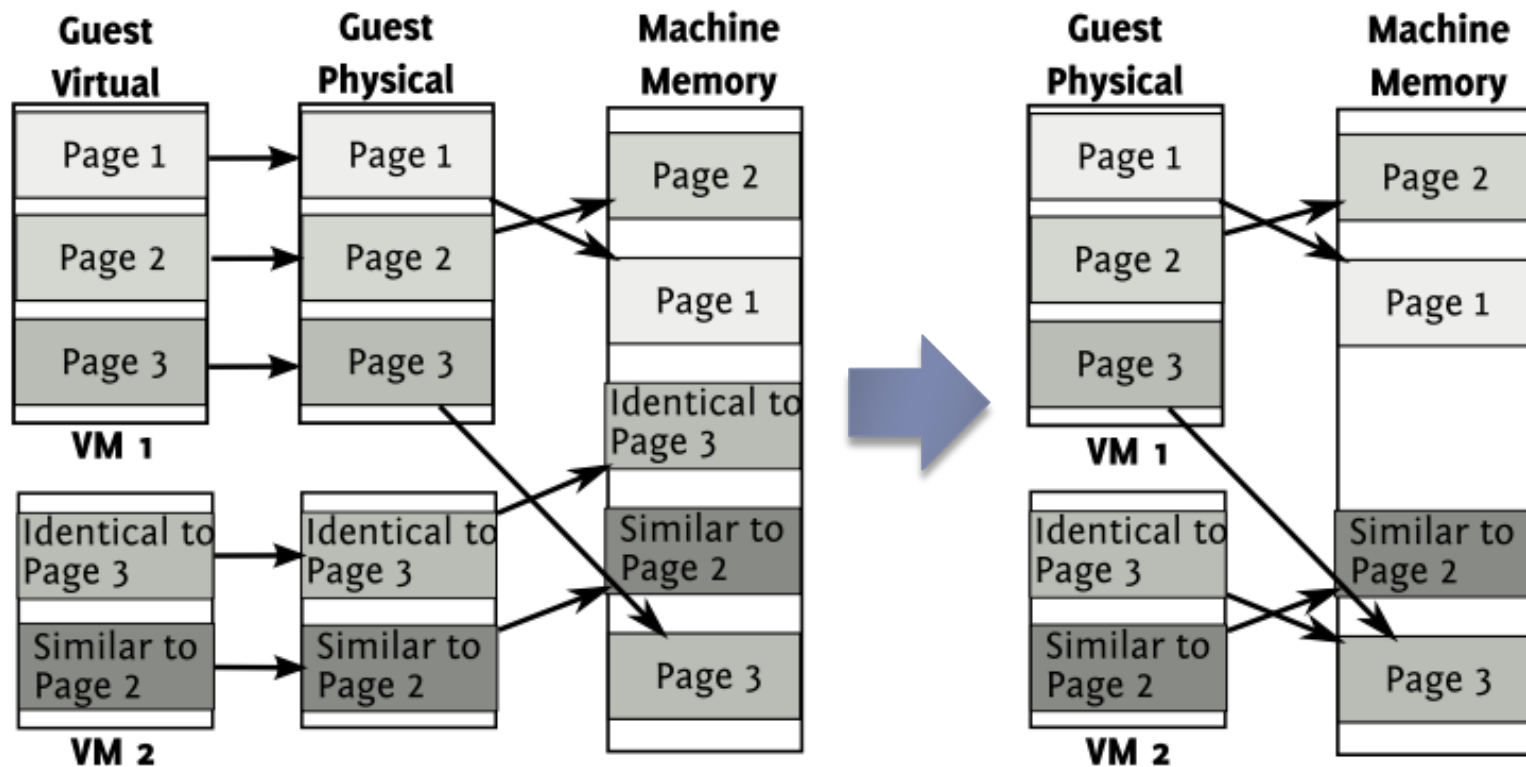
And a Bonus: Paging Machine Memory

- ▶ What if all VMs require all their allocated memory?
- ▶ And this allocated memory exceeds physical memory?
- ▶ Employ paging mechanism
 - ▶ Writing pages out to disk
- ▶ Use mechanism of locating infrequently used pages
- ▶ It is slow! So this operation must be infrequent!



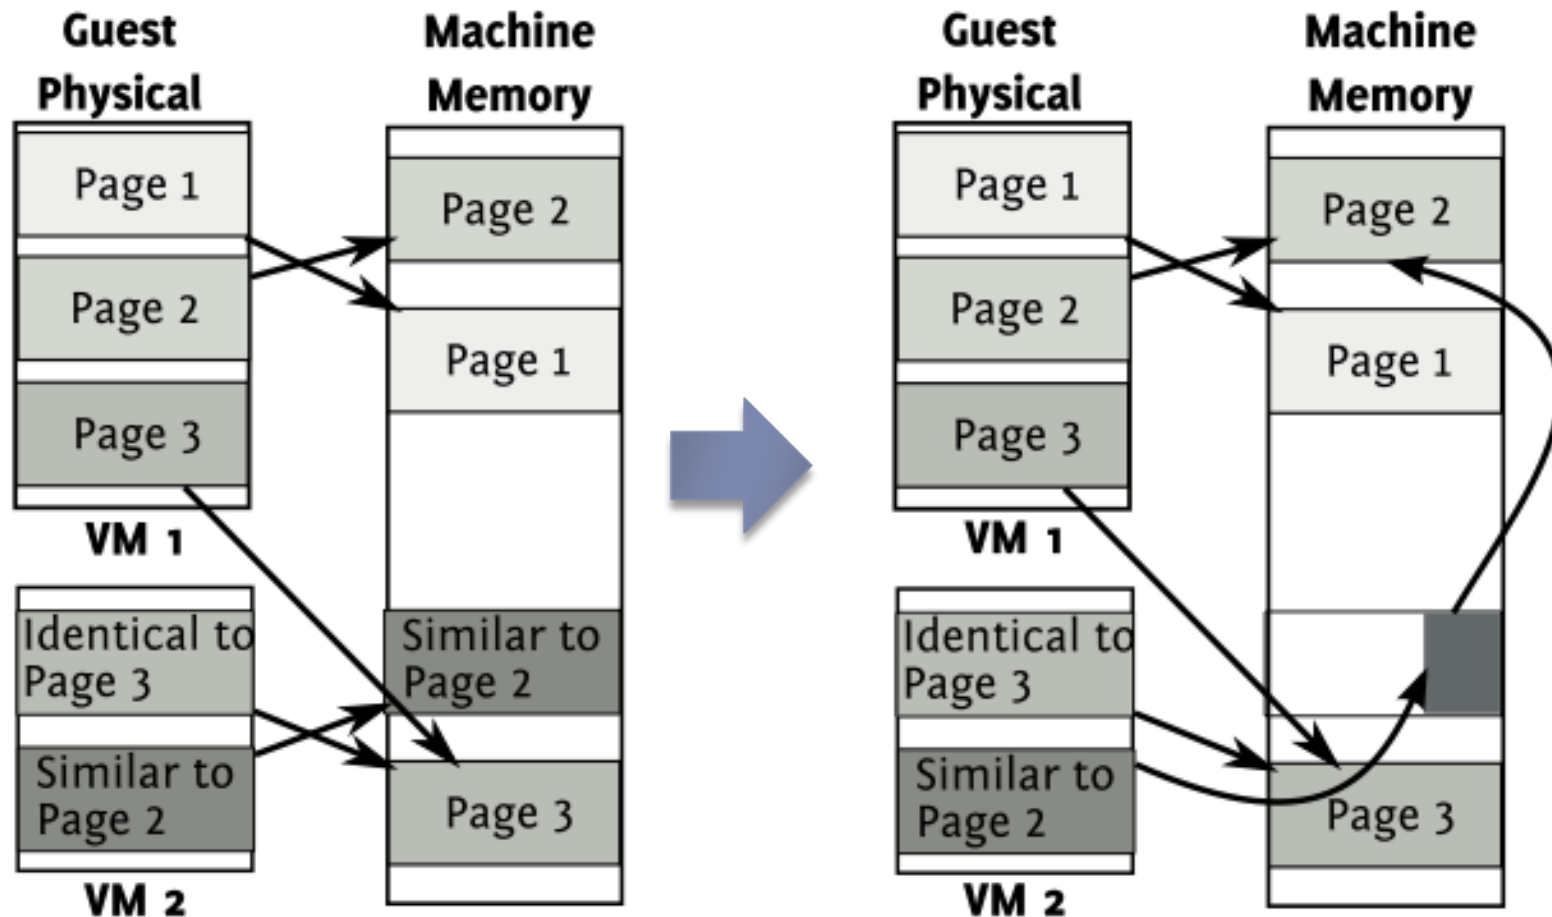
Step-by-step example

► Step 1: Page sharing



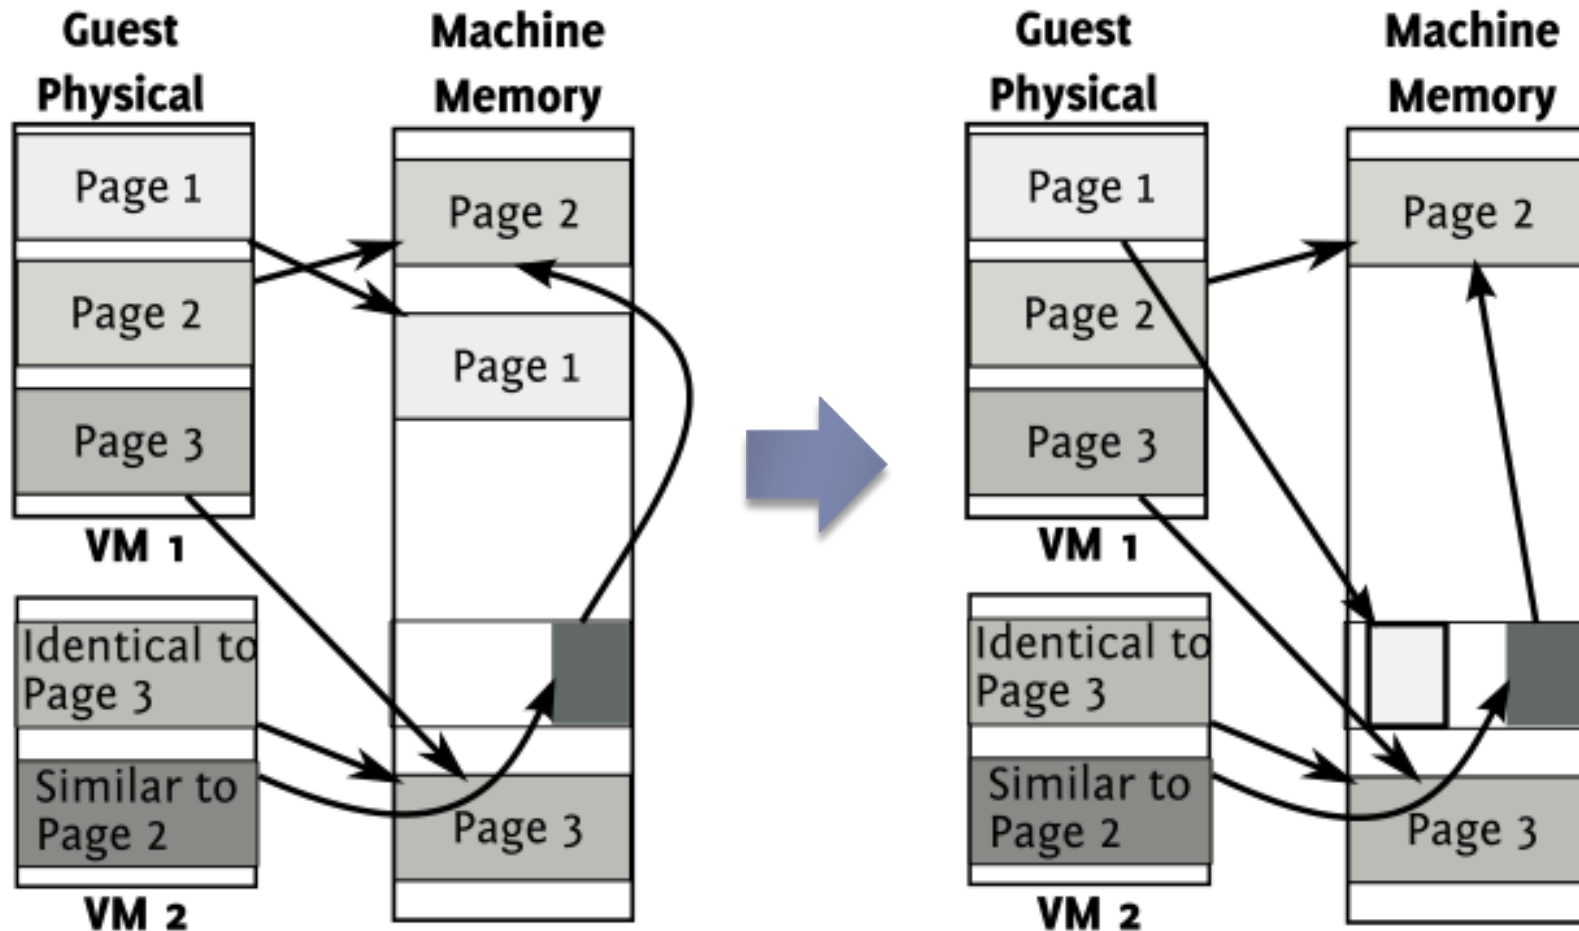
Step-by-step example

► Step 2: Patching




Step-by-step example

► Step 3: Compression



A few implementation notes

- ▶ Implementation on the top of Xen 3.0.4
- ▶ Roughly 14,500 lines of code
- ▶ Additional 20,000 lines from ports of existing algorithms
- ▶ And I'm not going to bore you with more details :)



Source: <http://barclay1720.tripod.com/>

Evaluation

- ▶ Evaluation setups:

- ▶ Homogeneous setup

- ▶ Mixed-1

- ▶ Windows XP SPI hosting RUBiS

- ▶ Debian 3.1 compiling Linux kernel

- ▶ Slackware 10.2 compiling Vim 7.0 followed by run of the 1mbench

- ▶ Mixed-2

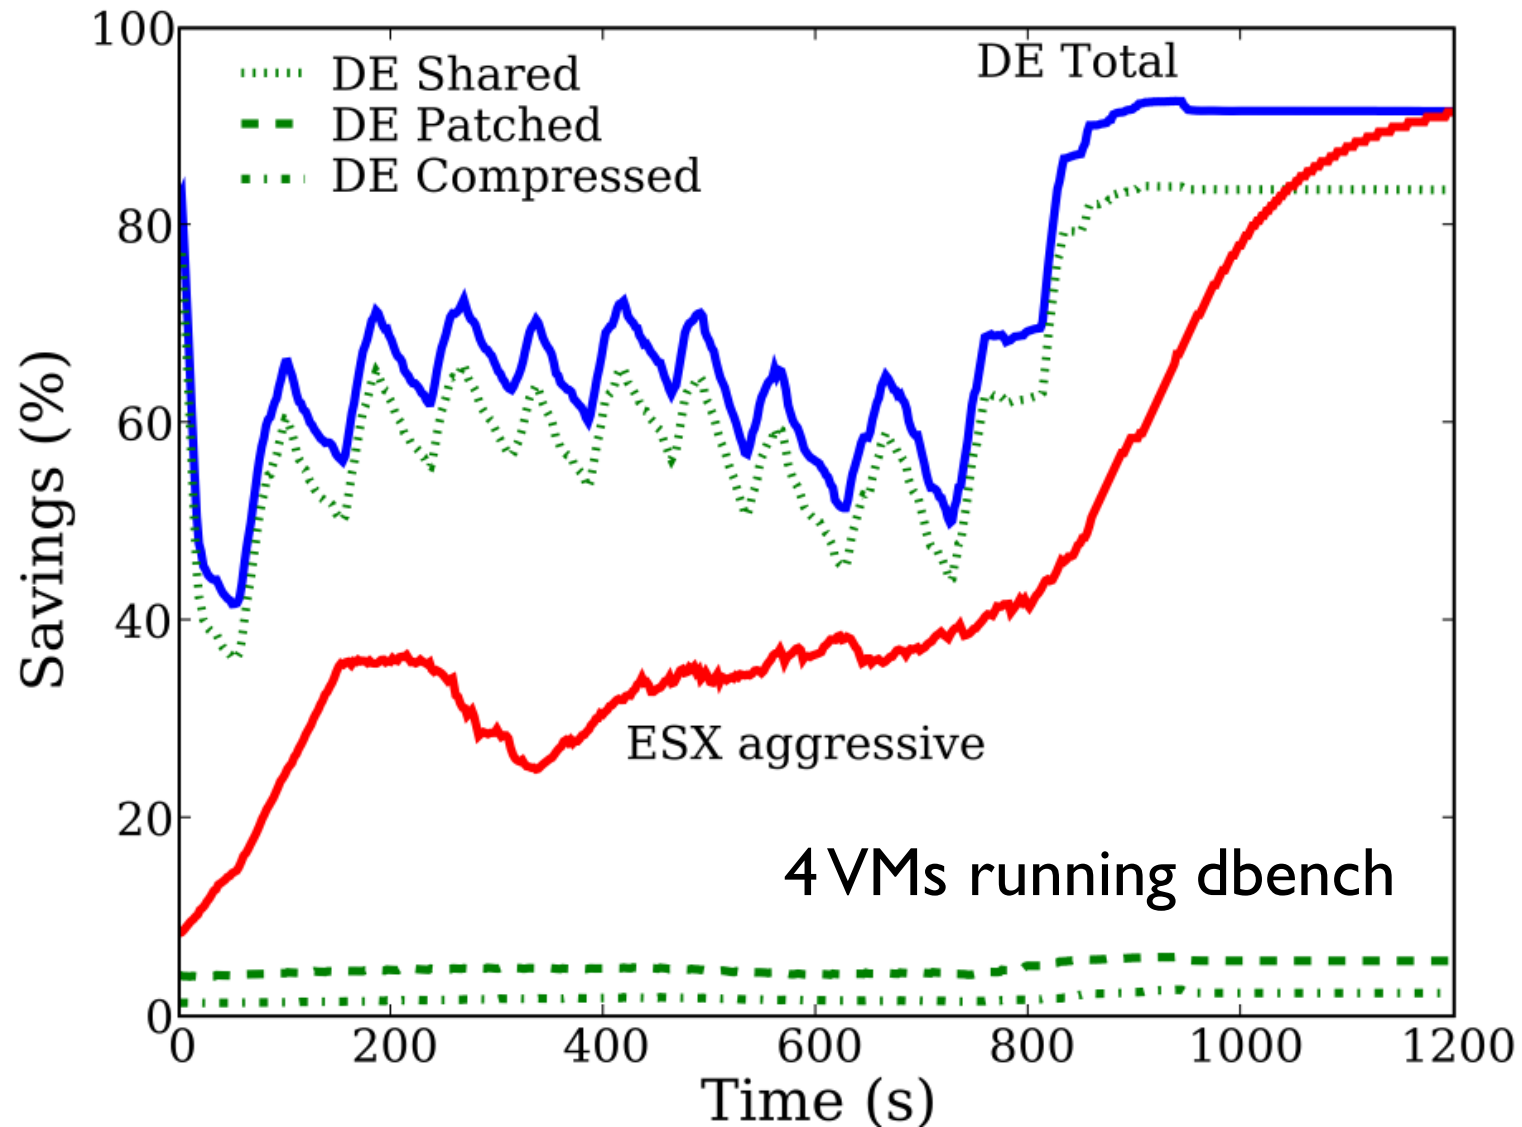
- ▶ Windows XP SPI running Apache 2.2.8 hosting a lot of static content (httperf running on a separate machine requesting these pages)

- ▶ Debian 3.1 running SysBench database benchmark


- ▶ Slackware 10.2 running dbench followed by IOZone




Memory savings for homogeneous setup



Memory savings for Mixed-1



Memory savings for Mixed-2



Performance evaluation

- ▶ Mixed-I setup


	Kernel Compile (sec)	Vim compile, lmbench (sec)	RUBiS requests	RUBiS response time(ms)
Baseline	670	620	3149	1280
DE	710	702	3130	1268

- ▶ Observed performance is within 7% of the baseline
- ▶ VMware ESX Server is within 5% of the baseline



How use reclaimed memory?

- ▶ We can spawn more virtual machines



Conclusions

- ▶ Harvesting identical pages across virtual machines
 - ▶ Works well on homogeneous systems
- ▶ Patching and in-memory page compression
 - ▶ Some improvement on heterogeneous systems
- ▶ Small performance overhead
- ▶ Article addresses many technical challenges
 - ▶ Algorithms
 - ▶ Xen limitations
 - ▶ Paging support
 - ▶ Clock mechanism for infrequently used pages location



Q & A



Source: <http://www.burrellesluce.com/>

